

JUSTIN CLARKE

Dubai, UAE • justinsavioclarke@outlook.com • +971 58537 6967
[linkedin.com/in/justinsavioclarke](https://www.linkedin.com/in/justinsavioclarke) • justinclarke.github.io

UAE Family Residence Visa • sponsorship not required

Analytics Engineer, with experience building ETL pipelines and tested data models across POS, ERP and e-commerce sources. Hands-on with Microsoft Fabric (Lakehouse, PySpark notebooks, pipelines), dbt, SQL, Python and incremental loading; built a 60-model dbt pipeline with 424 CI-enforced tests outside of work.

Experience

Analytics Engineer

VNS Solutions • Jan 2024 - Dec 2025

- Designed and maintained ETL pipelines ingesting POS, ERP and e-commerce data into a Microsoft Fabric Lakehouse using SQL transformations, Python notebooks and Fabric pipelines, producing reusable analytics datasets and improving reporting reliability.
- Built and maintained Microsoft Fabric data workflows using Lakehouse storage, PySpark transformation notebooks and orchestrated pipelines, creating curated analytics datasets for downstream BI consumption.
- Cut dataset refresh time from ~45 minutes to ~20 minutes on a multi-million-record reporting workflow by moving transformations from Power Query into SQL and implementing incremental loading.
- Implemented source-to-target validation and reconciliation checks across ETL workflows, detecting data-quality issues early and improving dataset reliability.

Product & Analytics Developer

LiteStore • Apr 2021 - Apr 2023

- GA4 instrumentation and A/B experiments drove a 20% conversion uplift across 11 retail brands
- Owned analytics, client reporting, and product direction as the sole technical contributor

Frontend & Brand Developer

Drop • Jan 2021 - Apr 2021

- Built React surfaces for a smart-lockbox delivery product; owned backend services, CD pipeline, hosting, and brand identity as the sole technical and creative contributor.

Featured Project

Off The Pace • Python • SQL • dbt • XGBoost • justinclarke.github.io/f1

- Engineered a medallion data pipeline: FastF1/Jolpica ingestion → Hive-partitioned Parquet → 60-model dbt transform layer on DuckDB, guarded by 424 CI-enforced data-quality tests.
- Also built a ChromaDB-backed knowledge-graph/lineage agent over a 700-file codebase for "what connects to what" queries agentic AI/LLM tooling on top of the pipeline.

Education

MBA in Business Analytics • BITS Pilani, Dubai • Feb 2026-2028 (*evening classes*)

MSc in Computer Science • Queen Mary, University of London • 2022-2023

BTech in Computer Science & Engineering • GITAM • 2018-2022

Technical Skills

- **Data Engineering:** Python, SQL, PL/SQL, ETL, dbt, PySpark, Microsoft Fabric (Lakehouse), DuckDB, Hive partitioning, dimensional modelling, data quality/reconciliation
- **Infra & Other:** Power BI, DAX, PostgreSQL, AWS, Docker, REST APIs, CI/CD, GitHub, ChromaDB/RAG